Evidence-led candidate matching

PANEL FLOW

The problem

Recruiters need a shortlist they can inspect, with evidence behind every requirement.

The product

Resume Intelligence combines hybrid retrieval, structured skill-gap analysis, and source-linked explanations.

The guardrail

Scores support recruiter review; they are not automated hiring decisions.

SAY

Show the login / workspace.

Explain the decision-support goal.

Set expectation: traceable matching, human decision.

One pipeline from document to evidence

PANEL FLOW

Prepare

Parse PDF, DOCX, or TXT; structure candidate facts; chunk and embed resume text.

Find

Understand job requirements, retrieve candidates with lexical and semantic signals, then rerank.

Explain

Matching, skill-gap, and qualification stages produce statuses and evidence for each result.

POINT TO

- Resume and job brief ingestion
- PostgreSQL / pgvector plus evidence records
- Fallback chains for unavailable AI providers
- Full diagram: architecture-diagram.pdf

A recruiter flow in the workspace

PANEL FLOW

Create a brief

Paste requirements or upload a PDF, DOCX, or TXT file.

Choose a search

Run a Top 3 or Top 5 shortlist with a saved brief or natural-language query.

Inspect the result

Open a candidate, compare profiles, and review matched, partial, and missing requirements.

DEMO

Open Jobs and show upload input.

Open Search and select a brief.

Keep one known candidate for the evidence walkthrough.

Live shortlist: follow one candidate

PANEL FLOW

Start the query

Use a prepared brief with a clear mandatory skill and experience requirement.

Read the shortlist

Show score breakdown, candidate rank, and retrieval/fallback status.

Verify the explanation

Open matched and partial skills; inspect quoted source text and page/source file.

Review risk

Show missing requirements and any validation flag; explain that absent evidence is not proof of absence.

DEMO CHECKPOINTS

Results view: top candidates

Candidate detail: evidence and page

Compare: side-by-side requirements

Mention education time is excluded from experience totals.

Quality is measured against human labels

PANEL FLOW

Build ground truth

Label query/candidate relevance and optional expected skill states: matched, partial, missing.

Run the same pipeline

Evaluation executes the live search path; unlabelled queries are reported as skipped.

Read the metrics

Ranking: precision@K, recall@K, MRR, NDCG. Skill gaps: accuracy and per-class precision/recall/F1.

Check grounding

Evidence claims are compared to persisted source evidence; report support rate and flagged results.

BE TRANSPARENT

Metrics reflect only supplied labels and evaluated outputs.

No labels means no fabricated accuracy.

Use a representative labeled dataset for meaningful metrics.

Tradeoffs, guardrails, and limits

PANEL FLOW

Retrieval tradeoff

Hybrid lexical + semantic improves recall; reranking costs additional time and compute.

Reliability

Provider chains fall back to local/deterministic paths; UI and API expose fallback status.

Traceability

Store evidence links and validation notes with each result for recruiter review.

Known limits

Gateway embedding compatibility needs confirmation; evaluate on your own consented corpus.

DESIGN CHOICES

Recall first, then rerank a bounded pool.

Keep source evidence alongside extracted fields.

Human review stays in the loop.

Close and invite questions

PANEL FLOW

Recap

The system turns resumes and role context into a ranked, evidence-linked shortlist.

Show next

Open the architecture PDF and README setup/sample-use sections.

Q&A

Reserve the final two minutes for panel questions.

Q&A | 08:00-10:00

How do you evaluate match quality?

What happens when an AI provider is unavailable?

How do recruiters verify an explanation?